




(Logistic Regression)

Aim: Implement Logistic Regression on a given Dataset with binary and multiclass labels.

Theory:

Logistic Regression is a statistical approach and a Machine Learning algorithm that is used for classification problems and is based on the concept of probability. It is used when the dependent variable (target) is categorical. It is widely used when the classification problem at hand is binary; true or false, yes or no, etc. For example, it can be used to predict whether an email is spam (1) or not (0). Logistics regression uses the sigmoid function to return the probability of a label.

Sigmoid Function is a mathematical function used to map the predicted values to probabilities. The function has the ability to map any real value into another value within a range of 0 and 1.



The rule is that the value of the logistic regression must be between 0 and 1. Due to the limitations of it not being able to go beyond the value 1, on a graph it forms a curve in the form of an "S". This is an easy way to identify the Sigmoid function or the logistic function. In regards to Logistic Regression, the concept



Department of Computer Science and Engineering (Data Science)

used is the threshold value. The threshold values help to define the probability of either 0 or 1. For example, values above the threshold value tend to 1, and a value below the threshold value tends to 0.

Type of Logistic Regression

1. Binomial: This means that there can be only two possible types of the dependent variables, such as 0 or 1, Yes or No, etc.
2. Multinomial: This means that there can be 3 or more possible unordered types of the dependent variable, such as "cat", "dogs", or "sheep"
3. Ordinal: This means that there can be 3 or more possible ordered types of dependent variables, such as "low", "Medium", or "High".

Binary Logistic Regression Major Assumptions

1. The dependent variable should be dichotomous in nature (e.g., presence vs. absent).
2. There should be no outliers in the data, which can be assessed by converting the continuous predictors to standardized scores, and removing values below -3.29 or greater than 3.29.
3. There should be no high correlations (multicollinearity) among the predictors. This can be assessed by a correlation matrix among the predictors. Tabachnick and Fidell (2013) suggest that as long correlation coefficients among independent variables are less than 0.90 the assumption is met. The aim of training the logistic regression model is to figure out the best weights for our linear model within the logistic regression. In machine learning, we compute the optimal weights by optimizing the cost function. **Cost function:** The cost function $J(\theta)$ is a formal representation of an objective that the algorithm is trying to achieve. In the case of logistic regression, the cost function is called LogLoss (or Cross-Entropy) and the goal is to minimize the following cost function equation:

4.
$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))]$$



Department of Computer Science and Engineering (Data Science)

Gradient descent is a method of changing weights based on the loss function for each data point. We calculate the LogLoss cost function at each input-output data point. We take a partial derivative of the weight and bias to get the slope of the cost function at each point. (No need to brush up on linear algebra and calculus right now. There are several matrix optimizations built into the Python library and Scikit-learn, which allow data science enthusiasts to unlock the power of advanced artificial intelligence without coding the answers themselves). Based on the slope, gradient descent updates the values for the bias and the set of weights, then reiterates the training loop over new values (moving a step closer to the desired goal). This iterative approach is repeated until a minimum error is reached, and gradient descent cannot minimize the cost function any further. We can change the speed at which we reach the optimal minimum by adjusting the learning rate. A high learning rate changes the weights more drastically, while a low learning rate changes them more slowly.

Lab Assignments to complete in this session:

Use the given dataset and perform the following tasks:

Dataset 1: Synthetic Dataset

Dataset 2: IRIS.csv

Dataset 3: Airlines_Passanger.csv

1. Perform required Logistic Regression from scratch on Dataset 1. Compare the F1 score of the LR model built from scratch and built using python library.
2. Perform Multimodal classification on Dataset 2 using python library.
3. Compare the results of Logistic Regression model with and without regularization.

Q1)

```
✓ [1] import numpy as np
0s      import pandas as pd
      import matplotlib.pyplot as plt
      from numpy import log,dot,exp,shape

✓ [2] from sklearn.datasets import make_classification
2s

✓ [3] X,y = make_classification()
0s      from sklearn.model_selection import train_test_split
      X_tr,X_te,y_tr,y_te = train_test_split(X,y,test_size=0.3)

✓ [4] def standardize(X_tr):
0s      for i in range(shape(X_tr)[1]):
          X_tr[:,i] = (X_tr[:,i] - np.mean(X_tr[:,i]))/np.std(X_tr[:,i])

✓ [5] def F1_score(y,y_hat):
0s      tp,tn,fp,fn = 0,0,0,0
      for i in range(len(y)):
          if y[i] == 1 and y_hat[i] == 1:
              tp += 1
          elif y[i] == 1 and y_hat[i] == 0:
              fn += 1
          elif y[i] == 0 and y_hat[i] == 1:
              fp += 1
          elif y[i] == 0 and y_hat[i] == 0:
              tn += 1
      precision = tp/(tp+fp)
      recall = tp/(tp+fn)
      f1_score = 2*precision*recall/(precision+recall)
      return f1_score
```

```
✓ [6] class LogidticRegression:
0s      def sigmoid(self,z):
          sig = 1/(1+exp(-z))
          return sig
      def initialize(self,X):
          weights = np.zeros((shape(X)[1]+1,1))
          X = np.c_[np.ones((shape(X)[0],1)),X]
          return weights,X
      def fit(self,X,y,alpha=0.001,iter=400):
          weights,X = self.initialize(X)
          def cost(theta):
              z = dot(X,theta)
              cost0 = y.T.dot(log(self.sigmoid(z)))
              cost1 = (1-y).T.dot(log(1-self.sigmoid(z)))
              cost = -((cost1 + cost0))/len(y)
              return cost
          cost_list = np.zeros(iter,)
          for i in range(iter):
              weights = weights - alpha*dot(X.T,self.sigmoid(dot(X,weights))-np.reshape(y,(len(y),1)))
              cost_list[i] = cost(weights)
          self.weights = weights
          return cost_list
      def predict(self,X):
          z = dot(self.initialize(X)[1],self.weights)
          lis = []
          for i in self.sigmoid(z):
              if i>0.5:
                  lis.append(1)
              else:
                  lis.append(0)
          return lis
```

```
✓ [7] standardize(X_tr)
0s standardize(X_te)
obj1 = LogitRegression()
model= obj1.fit(X_tr,y_tr)
y_pred = obj1.predict(X_te)
y_train = obj1.predict(X_tr)
#Let's see the f1-score for training and testing data
f1_score_tr = F1_score(y_tr,y_train)
f1_score_te = F1_score(y_te,y_pred)
print(f1_score_tr)
print(f1_score_te)
```


```
0.9315068493150684
0.7333333333333333
```

```
✓ [8] from sklearn.linear_model import LogisticRegression
0s from sklearn.metrics import f1_score
model = LogisticRegression().fit(X_tr,y_tr)
y_pred = model.predict(X_te)
f1_score(y_te,y_pred)
```

```
0.6875
```

```
✓ [9] def plot_gradient_descent(X, y, alpha=0.001, iter=400):
0s obj = LogitRegression()
cost_list = obj.fit(X, y, alpha=alpha, iter=iter)
plt.plot(np.arange(iter), cost_list)
plt.xlabel('Iterations')
plt.ylabel('Cost')
plt.title('Gradient Descent')
plt.show()
```

```
✓ [10] plot_gradient_descent(X_tr, y_tr, alpha=0.01, iter=500)
0s
```



```


✓ [11] def plot_training_testing_curve(X_tr, y_tr, X_te, y_te, alpha=0.001, iter=400):
0s
    obj = LogitRegression()
    train_scores = []
    test_scores = []
    for i in range(1, iter+1):
        obj.fit(X_tr, y_tr, alpha=alpha, iter=i)
        y_train = obj.predict(X_tr)
        y_test = obj.predict(X_te)
        train_score = F1_score(y_tr, y_train)
        test_score = F1_score(y_te, y_test)
        train_scores.append(train_score)
        test_scores.append(test_score)
    plt.plot(np.arange(1, iter+1), train_scores, label='Training')
    plt.plot(np.arange(1, iter+1), test_scores, label='Testing')
    plt.xlabel('Iterations')
    plt.ylabel('F1-Score')
    plt.title('Training vs Testing Curve')
    plt.legend()
    plt.show()

```

```

✓ 21s plot_training_testing_curve(X_tr, y_tr, X_te, y_te, alpha=0.01, iter=500)

```



Q2)

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
[346] df=pd.read_csv('/content/train.csv')
```

```
df.head()
```



	Unnamed: 0	id	Gender	Customer Type	Age	Type of Travel	Class	Flight Distance	Inflight wifi service	Departure/Arrival time convenient	...	Inflight entertainment	On-board service	Leg room service	Baggage handling	Checkin service
0	0	70172	Male	Loyal Customer	13	Personal Travel	Eco Plus	460	3	4	...	5	4	3	4	
1	1	5047	Male	disloyal Customer	25	Business travel	Business	235	3	2	...	1	1	5	3	
2	2	110028	Female	Loyal Customer	26	Business travel	Business	1142	2	2	...	5	4	3	4	
3	3	24026	Female	Loyal Customer	25	Business travel	Business	562	2	5	...	2	2	5	3	
4	4	119299	Male	Loyal Customer	61	Business travel	Business	214	3	3	...	3	3	4	4	

5 rows × 25 columns



```
[348] df.drop(columns='Unnamed: 0',inplace=True)
```

```
[349] df.drop(columns='id',inplace=True)
```

```
df.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 103904 entries, 0 to 103903
Data columns (total 23 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Gender                                103904 non-null object
1   Customer Type                         103904 non-null object
2   Age                                   103904 non-null int64
3   Type of Travel                        103904 non-null object
4   Class                                 103904 non-null object
5   Flight Distance                       103904 non-null int64
6   Inflight wifi service                 103904 non-null int64
7   Departure/Arrival time convenient    103904 non-null int64
8   Ease of Online booking               103904 non-null int64
9   Gate location                        103904 non-null int64
10  Food and drink                       103904 non-null int64
11  Online boarding                      103904 non-null int64
12  Seat comfort                         103904 non-null int64
13  Inflight entertainment               103904 non-null int64
14  On-board service                     103904 non-null int64
15  Leg room service                     103904 non-null int64
16  Baggage handling                     103904 non-null int64
17  Checkin service                      103904 non-null int64
18  Inflight service                     103904 non-null int64
19  Cleanliness                          103904 non-null int64
20  Departure Delay in Minutes           103904 non-null int64
```

```
✓ [351] from sklearn.preprocessing import LabelEncoder
0s labelencoder=LabelEncoder()
df['satisfaction']=labelencoder.fit_transform(df['satisfaction'])
```

```
✓ [352] df['satisfaction'].astype(int)
0s
      0      0
      1      0
      2      1
      3      0
      4      1
      ..
103899    0
103900    1
103901    0
103902    0
103903    0
Name: satisfaction, Length: 103904, dtype: int64
```

```
✓ [353] def find_categorical_columns(df):
0s     cat_cols = []
     for col in df.columns:
         if pd.api.types.is_categorical_dtype(df[col]):
             cat_cols.append(col)
         elif pd.api.types.is_object_dtype(df[col]):
             cat_cols.append(col)
     return cat_cols
```

```
✓ [354] cat_columns=find_categorical_columns(df)
0s
```

```
✓ [355] def encode_categorical(df, cat_cols):
0s     df=pd.get_dummies(df, columns=cat_cols, prefix=cat_cols, prefix_sep='_')
     return(df)
```

```
✓ [356] cat_cols
0s
['Gender', 'Customer Type', 'Type of Travel', 'Class']
```

```
✓ [357] cat_cols
0s
['Gender', 'Customer Type', 'Type of Travel', 'Class']
```

```
✓ [358] df=encode_categorical(df,cat_cols)
0s
```

```
✓ [361] df.dropna(inplace=True)
0s
```

```
✓ [362] x=df.drop(columns='satisfaction').values
0s     y=df['satisfaction'].values
```

```
✓ [363] from sklearn import linear_model
8s model=linear_model.LogisticRegression(C=1e40, solver='newton-cg')
model.fit(x,y)
```

```
LogisticRegression
LogisticRegression(C=1e+40, solver='newton-cg')
```

```
✓ [364] lr=linear_model.LogisticRegression(penalty='l2', C=1.0, solver='lbfgs', max_iter=1000)
0s
```

```
✓ 29s ▶ lr.fit(x,y)
```

⚠ /usr/local/lib/python3.9/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
LogisticRegression
LogisticRegression(max_iter=1000)
```

```
✓ [366] df1=pd.read_csv('/content/test.csv')
0s
```



```
✓ [366] df1=pd.read_csv('/content/test.csv')
```

```
✓ [368] df1.dropna(inplace=True)
```

```
✓ [369] df1.drop(columns='Unnamed: 0',inplace=True)
```

```
✓ [370] df1['satisfaction']=labelencoder.fit_transform(df1['satisfaction'])
```

```
✓ [371] cat_cols=find_categorical_columns(df1)
```

```
✓ [372] df1=encode_categorical(df1,cat_cols)
```

```
✓ [373] df1.head()
```

	id	Age	Flight Distance	Inflight wifi service	Departure/Arrival time convenient	Ease of Online booking	Gate location	Food and drink	Online boarding	Seat comfort	...	satisfaction	Gender_Female	Gender_Male	Customer Type_Loyal Customer	Customer Type_disloyal Customer
0	19556	52	160	5	4	3	4	3	4	3	...	1	1	0	1	0
1	90035	36	2863	1	1	3	1	5	4	5	...	1	1	0	1	0
2	12360	20	192	2	0	2	4	2	2	2	...	0	0	1	0	1
3	77959	44	3377	0	0	0	2	3	4	4	...	1	0	1	1	0
4	36875	49	1182	2	3	4	3	4	1	2	...	1	1	0	1	0

5 rows × 29 columns



```
✓ [374] df1.drop(columns='id',inplace=True)
```

```
✓ [375] x=df1.drop(columns='satisfaction')
y=df1['satisfaction']
```

Q3)

Accuracy without regularisation:

```
✓ [376] y_pred=model.predict(x)
```

```
/usr/local/lib/python3.9/dist-packages/sklearn/base.py:432: UserWarning: X has feature names, but LogisticRegression was fitted without feature names
warnings.warn(
```

```
✓ [377] from sklearn.metrics import accuracy_score
accuracy_score(y,y_pred)
```

```
0.8717027768122658
```

Accuracy with Lasso regularisation:

```
✓ [378] y_pred=lr.predict(x)
```

```
/usr/local/lib/python3.9/dist-packages/sklearn/base.py:432: UserWarning: X has feature names, but LogisticRegression was fitted without feature names
warnings.warn(
```

```
✓ [379] accuracy_score(y,y_pred)
```

```
0.8688062410690148
```

Conclusion: Logistic regression has successfully been performed with & without Lasso regularisation.