



**Department of Computer Science and Engineering (Data Science)**  
High Performance Computing Laboratory (DJ19DSL802)

**HPC Experiment-5**

*Bhuvi Ghosh*  
*60009210191*

Aim: Accelerating a For Loop with Multiple Blocks of Threads

Theory:

There is a limit to the number of threads that can exist in a thread block: 1024 to be precise. In order to increase the amount of parallelism in accelerated applications, we must be able to coordinate among multiple thread blocks.

CUDA Kernels have access to a special variable that gives the number of threads in a block: `blockDim.x`. Using this variable, in conjunction with `blockIdx.x` and `threadIdx.x`, increased parallelization can be accomplished by organizing parallel execution across multiple blocks of multiple threads with the idiomatic expression `threadIdx.x + blockIdx.x * blockDim.x`. Here is a detailed example.

The execution configuration `<<<10, 10>>>` would launch a grid with a total of 100 threads, contained in 10 blocks of 10 threads. We would therefore hope for each thread to have the ability to calculate some index unique to itself between 0 and 99.

1. If block `blockIdx.x` equals 0, then `blockIdx.x * blockDim.x` is 0. Adding to 0 the possible `threadIdx.x` values 0 through 9, then we can generate the indices 0 through 9 within the 100 thread grid.
2. If block `blockIdx.x` equals 1, then `blockIdx.x * blockDim.x` is 10. Adding to 10 the possible `threadIdx.x` values 0 through 9, then we can generate the indices 10 through 19 within the 100 thread grid.
3. If block `blockIdx.x` equals 5, then `blockIdx.x * blockDim.x` is 50. Adding to 50 the possible `threadIdx.x` values 0 through 9, then we can generate the indices 50 through 59 within the 100 thread grid.
4. If block `blockIdx.x` equals 9, then `blockIdx.x * blockDim.x` is 90. Adding to 90 the possible `threadIdx.x` values 0 through 9, then we can generate the indices 90 through 99 within the 100 thread grid.

**Lab Experiment to be performed:**



**Department of Computer Science and Engineering (Data Science)**  
High Performance Computing Laboratory (DJ19DSL802)

Currently, the loop function inside 02-multi-block-loop.cu runs a for loop that will serially print the numbers 0 through 9. Refactor the loop function to be a CUDA kernel which will launch to execute N iterations in parallel. After successfully refactoring, the numbers 0 through 9 should still be printed. For this exercise, as an additional constraint, use an execution configuration that launches at least 2 blocks of threads. Refer to the solution if you get stuck.

```
[11] !nvcc -arch=sm_70 -o multi-block-loop /content/02-multi-block-loop.cu -run
0
1
2
3
4
5
6
7
8
9
```



**Department of Computer Science and Engineering (Data Science)**  
High Performance Computing Laboratory (DJ19DSL802)

02-multi-block-loop.cu X

```
1 #include <stdio.h>
2
3 /*
4  * Refactor `loop` to be a CUDA Kernel. The new kernel should
5  * only do the work of 1 iteration of the original loop.
6  */
7
8 __global__ void loop()
9 {
10     /*
11     * This idiomatic expression gives each thread
12     * a unique index within the entire grid.
13     */
14     int i = blockIdx.x * blockDim.x + threadIdx.x;
15     printf("%d\n", i);
16 }
17 }
```

```
18
19 int main()
20 {
21     /*
22     * Additional execution configurations that would
23     * work and meet the exercises constraints are:
24     *
25     * <<<5, 2>>>
26     * <<<10, 1>>>
27     */
28     loop<<<2, 5>>>();
29     cudaDeviceSynchronize();
30 }
31 }
32 }
```